

Case Challenge: Smart Financial Coach

FinAI

By Anahita Dinesh

1. Project Overview

FinAI is a lang-chain based full-stack application designed to bridge the gap between raw financial data and actionable behavioral change. The application allows users to ask questions in plain English and receive automated, data-driven insights.

Github: <https://github.com/annahita20/personal-finance-insights-langchain>

Demo: <https://vimeo.com/1160883209?share=copy&fl=sv&fe=ci>

2. Technical Architecture

The system follows a decoupled architecture, separating data persistence, the AI orchestration layer, and the visualization.

2.1 Tech Stack

- Frontend: React.js with Tailwind CSS for a responsive UI and Recharts for dynamic data visualization.
- Backend: Flask (Python) serving as the API layer.
- AI Orchestration: LangChain for managing LLM prompts, entity extraction, and SQL generation.
- LLM: Google Gemini (via API) for natural language processing.
- Database: SQLite for local, lightweight relational data storage.


2.2 Functionality

The application handles two primary workflows: the AI Chat and the Automated Insights Dashboard.

AI Chat Pipeline

- Entity Extraction: The LLM parses the user's natural language question to identify key entities (ex, coffee, last month, Starbucks).
- Schema Mapping: The system matches extracted entities against the database schema and unique column values.
- SQL Translation: LangChain prompts the LLM to generate a syntactically correct SQLite query based on the user's intent.
- Data Execution: The query is executed against the SQLite database.
- Natural Language Synthesis: The raw data results are passed back to the LLM to be formatted into a conversational, "coaching" style response.


Finance SQL Pipeline



Dashboard & Insights Flow

1. Data Retrieval: The dashboard fetches aggregated transaction data.
2. Visualization: React-Recharts renders spending trends and category breakdowns.
3. Anomaly Detection: The AI-Insights engine scans the returned dataset for patterns, such as "gray charges" (forgotten subscriptions) or spending spikes, returning them as proactive alerts.


Finance Dashboard General Flow



3. Data Modeling

The system utilizes a relational schema to ensure strict financial integrity and high-performance aggregation.

1. ACCOUNTS: Stores account metadata and balances.
2. TRANSACTIONS: The core of the application revolves around the TRANSACTIONS table, which acts as a central ledger. It captures granular spend data. Includes other_dynamic_cols to support varied CSV imports from different banking institutions.
3. MONTHLY_BUDGETS: Defines spending limits per category for variance analysis.
4. GLOBAL_SEARCH_INDEX: A specialized index table used by the AI pipeline to map natural language terms to specific database records.



Design Decision: SQL vs. Graph

While Graph databases excel at relationship mapping, a Relational (SQL) approach was chosen for:

- ACID Compliance: Ensuring that balance updates and transaction logs are atomic and consistent.
- Aggregation Speed: SQL is natively optimized for calculating sums, averages, and month-over-month variances across thousands of rows.

4. Key Features & Implementation

1. Natural Language Querying (NLQ): Users can ask, "How much did I spend on dining out in January?" without needing to navigate complex filters.
2. Subscription Detection: The system identifies recurring transaction descriptions with consistent amounts to flag potential "gray charges."
3. Goal Forecasting: By analyzing historical burn rates against financial goals, the AI predicts the likelihood of reaching a savings target by a specific date.

5. Security & Trust

- Local Processing: The use of SQLite allows for local data persistence, minimizing the footprint of sensitive financial records.
- API Safety: The system utilizes environment variables (.env) for API key management, ensuring credentials are never hardcoded or exposed.

6. Future Enhancements

- Vector Embeddings: Implementing a vector store for semantic search of transaction descriptions to improve entity matching.
- Multi-User Support: Migrating from SQLite to PostgreSQL for multi-tenant capabilities.